

This is a companion notebook for the book [Deep Learning with Python, Second Edition](#). For readability, it only contains runnable code blocks and section titles, and omits everything else in the book: text paragraphs, figures, and pseudocode.

**If you want to be able to follow what's going on, I recommend reading the notebook side by side with your copy of the book.**

This notebook was generated for TensorFlow 2.6.

## ✓ Getting started with neural networks: Classification and regression

### ✓ Classifying movie reviews: A binary classification example

#### ✓ The IMDB dataset

##### Loading the IMDB dataset

```
from tensorflow.keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
```

```
train_data[0]
```

```
⇒ [1,
   14,
   22,
   16,
   43,
   530,
   973,
   1622,
   1385,
   65,
   458,
   4468,
   66,
   3941,
   4,
   173,
   36,
   256,
```

```
5,  
25,  
100,  
43,  
838,  
112,  
50,  
670,  
2,  
9,  
35,  
480,  
284,  
5,  
150,  
4,  
172,  
112,  
167,  
2,  
336,  
385,  
39,  
4,  
172,  
4536,  
1111,  
17,  
546,  
38,  
13,  
447,  
4,  
192,  
50,  
16,  
6,  
147,  
2025,  
19.
```

```
train_labels[0]
```

```
⇒ 1
```

```
max([max(sequence) for sequence in train_data])
```

```
⇒ 9999
```

## Decoding reviews back to text

```
word_index = imdb.get_word_index()  
reverse_word_index = dict(  

```

```
[(value, key) for (key, value) in word_index.items()])
decoded_review = " ".join(
    [reverse_word_index.get(i - 3, "?") for i in train_data[0]])
```

## ✓ Preparing the data

### Encoding the integer sequences via multi-hot encoding

```
import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
```

```
x_train[0]
```

```
→ array([0., 1., 1., ..., 0., 0., 0.])
```

```
y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

## ✓ Building your model

### Model definition

```
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

### Compiling the model

```
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

## ✓ Validating your approach

### Setting aside a validation set

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

### Training your model

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=4,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
➡ Epoch 1/4
30/30 ————— 3s 67ms/step - accuracy: 0.6808 - loss: 0.6103 - val_accuracy
Epoch 2/4
30/30 ————— 1s 47ms/step - accuracy: 0.8787 - loss: 0.3659 - val_accuracy
Epoch 3/4
30/30 ————— 2s 37ms/step - accuracy: 0.9166 - loss: 0.2607 - val_accuracy
Epoch 4/4
30/30 ————— 1s 34ms/step - accuracy: 0.9265 - loss: 0.2167 - val_accuracy
```


```
history_dict = history.history
history_dict.keys()
```

```
➡ dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

### Plotting the training and validation loss


```
import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
```

```
plt.xlabel("Epochs")  
plt.ylabel("Loss")  
plt.legend()  
plt.show()
```



## Plotting the training and validation accuracy

```
plt.clf()  
acc = history_dict["accuracy"]  
val_acc = history_dict["val_accuracy"]  
plt.plot(epochs, acc, "bo", label="Training acc")  
plt.plot(epochs, val_acc, "b", label="Validation acc")  
plt.title("Training and validation accuracy")  
plt.xlabel("Epochs")  
plt.ylabel("Accuracy")  
plt.legend()  
plt.show()
```



Training Results:

```
results = model.evaluate(x_test, y_test)
```



**782/782** ————— 2s 3ms/step - accuracy: 0.8782 - loss: 0.2941

Test Results

```
results
```

```
#creating a results table
```

```
results_history = {}
results_history["base_model"] = {
    "loss": results[0],
    "accuracy": results[1],
    "history": history.history
}
```

The training and test results from the base model are relatively close, at an accuracy of .86, and a loss of .32. For this 'baseline' model, I kept the 4 epochs as accuracy seemed to peak there and

taper off after 4. In the next few iterations, I will use the suggestions from the assignment to see which changes affect the model for the better.

## ✓ Question 1: Add a hidden layer to the model to see effects on test and training accuracy.

```
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"), # added hidden layer
    layers.Dense(1, activation="sigmoid")
])
```

```
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=4,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
history_dict = history.history
history_dict.keys()
```

```
↩ Epoch 1/4
30/30 ————— 4s 83ms/step - accuracy: 0.6827 - loss: 0.6214 - val_accuracy
Epoch 2/4
30/30 ————— 1s 36ms/step - accuracy: 0.8912 - loss: 0.3391 - val_accuracy
Epoch 3/4
30/30 ————— 1s 38ms/step - accuracy: 0.9198 - loss: 0.2353 - val_accuracy
Epoch 4/4
30/30 ————— 1s 36ms/step - accuracy: 0.9459 - loss: 0.1759 - val_accuracy
dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

Training Results:

```
results = model.evaluate(x_test, y_test)
```

↩ 782/782 ————— 2s 3ms/step - accuracy: 0.8791 - loss: 0.2977

Test Results:

```
results
```

```
results_history["model_with_added_layer"] = {
    "loss": results[0],
    "accuracy": results[1],
    "history": history.history
}
```

By adding a hidden layer, the model gets marginally better, at .87 for accuracy, and .31 loss for the training data set, and .3 for the test data set. I don't know that adding the extra layer, time, and computational capacity is worth it for these slightly better results. Next, I'll change the number of hidden units in each layer.

## ✓ Question 2: More or fewer hidden units per layer.

```
model = keras.Sequential([
    layers.Dense(8, activation="relu"),
    layers.Dense(8, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=4,
                    batch_size=512,
                    validation_data=(x_val, y_val))

history_dict = history.history
history_dict.keys()
```



```

Epoch 1/4
30/30 ————— 4s 82ms/step - accuracy: 0.6769 - loss: 0.6307 - val_accuracy
Epoch 2/4
30/30 ————— 1s 34ms/step - accuracy: 0.8871 - loss: 0.4084 - val_accuracy
Epoch 3/4
30/30 ————— 1s 32ms/step - accuracy: 0.9169 - loss: 0.3047 - val_accuracy
Epoch 4/4
30/30 ————— 1s 34ms/step - accuracy: 0.9329 - loss: 0.2362 - val_accuracy
dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])

```

```
results = model.evaluate(x_test, y_test)
```

```

782/782 ————— 2s 2ms/step - accuracy: 0.8831 - loss: 0.2998

```

```
results
```

```

results_history["model_with_changed_hidden_units"] = {
    "loss": results[0],
    "accuracy": results[1],
    "history": history.history
}

```

I landed on 8 hidden units per layer, because the results had the best balance of high accuracy and lower loss function vs 32, 64, or 128.

## Question 3: Using the MSE Loss function instead of binary\_crossentropy

```

model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

```

```

model.compile(optimizer="rmsprop",
              loss="mse",
              metrics=["accuracy"])

```

```

x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

```

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=4,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
history_dict = history.history
history_dict.keys()
```

```
Epoch 1/4
30/30 ————— 4s 78ms/step - accuracy: 0.6856 - loss: 0.2189 - val_accuracy
Epoch 2/4
30/30 ————— 1s 36ms/step - accuracy: 0.8762 - loss: 0.1213 - val_accuracy
Epoch 3/4
30/30 ————— 1s 37ms/step - accuracy: 0.9096 - loss: 0.0848 - val_accuracy
Epoch 4/4
30/30 ————— 1s 36ms/step - accuracy: 0.9223 - loss: 0.0700 - val_accuracy
dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

```
results = model.evaluate(x_test, y_test)
```

```
782/782 ————— 2s 2ms/step - accuracy: 0.8832 - loss: 0.0909
```

```
results
```

```
results_history["model_with_mse"] = {
    "loss": results[0],
    "accuracy": results[1],
    "history": history.history
}
```

## ✓ Question 4: Use tanh activation instead of relu.

```
model = keras.Sequential([
    layers.Dense(16, activation="tanh"),
    layers.Dense(16, activation="tanh"),
    layers.Dense(1, activation="sigmoid")
])
```

```
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

```

x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=4,
                    batch_size=512,
                    validation_data=(x_val, y_val))

```

```

history_dict = history.history
history_dict.keys()

```

```

Epoch 1/4
30/30 ————— 4s 93ms/step - accuracy: 0.7017 - loss: 0.5772 - val_accuracy
Epoch 2/4
30/30 ————— 3s 37ms/step - accuracy: 0.8998 - loss: 0.3152 - val_accuracy
Epoch 3/4
30/30 ————— 1s 33ms/step - accuracy: 0.9254 - loss: 0.2255 - val_accuracy
Epoch 4/4
30/30 ————— 1s 36ms/step - accuracy: 0.9442 - loss: 0.1724 - val_accuracy
dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])

```

```

results = model.evaluate(x_test, y_test)

```

```

782/782 ————— 2s 3ms/step - accuracy: 0.8656 - loss: 0.3309

```

```

results

```

```

results_history["model_with_tanh"] = {
    "loss": results[0],
    "accuracy": results[1],
    "history": history.history
}

```

Using the tahn function instead of relu slightly improved accuracy and loss, but still had issues with overfitting to the training data.

✓ Question 5: Adding something we learned in class. (I'm adding Dropout to reduce overfitting)

```
from tensorflow.keras.layers import Dense, Dropout
```

```
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    Dropout(0.3),
    layers.Dense(16, activation="relu"),
    Dropout(0.3),
    layers.Dense(1, activation="sigmoid")
])
```

```
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=4,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
history_dict = history.history
history_dict.keys()
```

```
Epoch 1/4
30/30 ————— 3s 66ms/step - accuracy: 0.6444 - loss: 0.6284 - val_accuracy
Epoch 2/4
30/30 ————— 2s 37ms/step - accuracy: 0.8429 - loss: 0.4137 - val_accuracy
Epoch 3/4
30/30 ————— 1s 37ms/step - accuracy: 0.8822 - loss: 0.3229 - val_accuracy
Epoch 4/4
30/30 ————— 1s 38ms/step - accuracy: 0.9066 - loss: 0.2671 - val_accuracy
dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

```
results = model.evaluate(x_test, y_test)
```

```
782/782 ————— 2s 3ms/step - accuracy: 0.8828 - loss: 0.2886
```

```
results
```

```
results_history["model_with_dropout"] = {
    "loss": results[0],
    "accuracy": results[1],
```

```
"history": history.history
}
```

Adding the dropout didn't quite work how I expected. Although it did improve overfitting from previous versions of the model, it didn't improve the overall accuracy or loss function of the model.

## ✓ Plotting the results in graphs as well as a table

```
# Plot accuracy and loss for all the different models
def plot_model_results(model_name, history):
    history_dict = history
    loss_values = history_dict["loss"]
    val_loss_values = history_dict["val_loss"]
    epochs = range(1, len(loss_values) + 1)



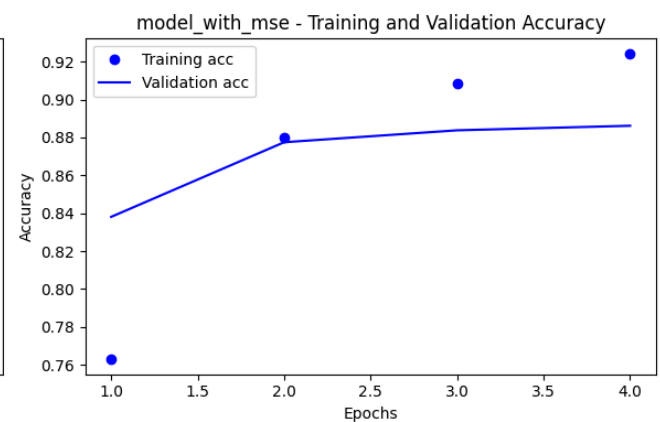
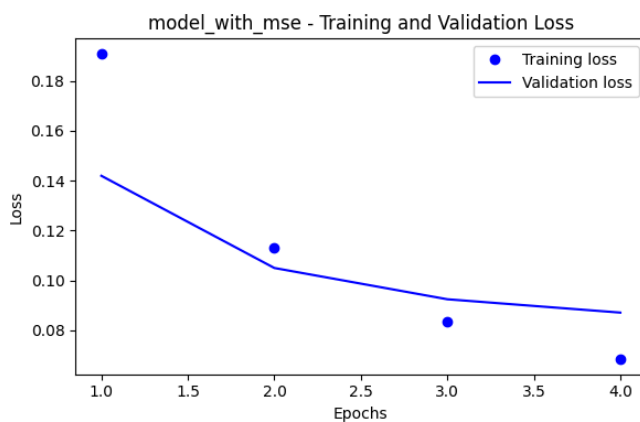



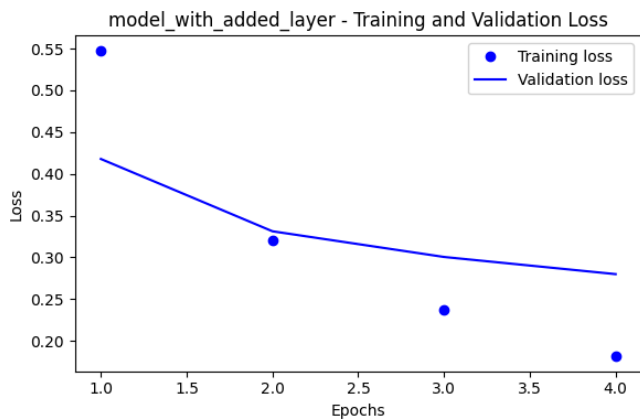


    plt.figure(figsize=(12, 4))

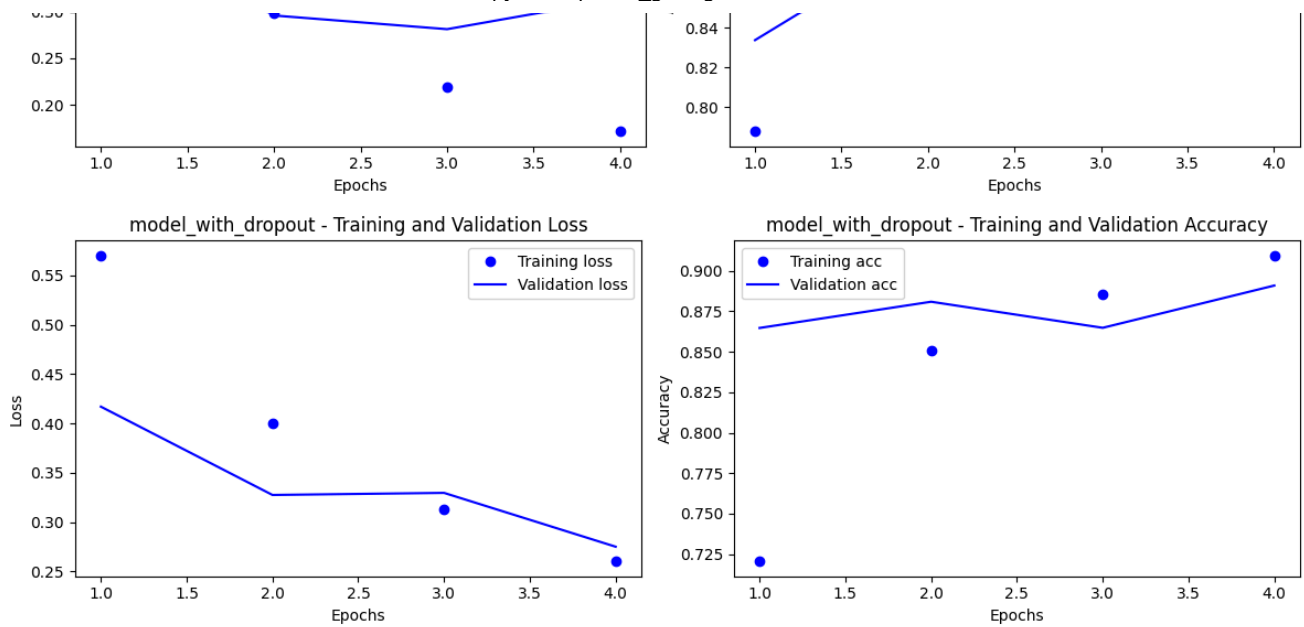
    plt.subplot(1, 2, 1)
    plt.plot(epochs, loss_values, "bo", label="Training loss")
    plt.plot(epochs, val_loss_values, "b", label="Validation loss")
    plt.title(f"{model_name} - Training and Validation Loss")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()

    plt.subplot(1, 2, 2)
    acc = history_dict["accuracy"]
    val_acc = history_dict["val_accuracy"]
    plt.plot(epochs, acc, "bo", label="Training acc")
    plt.plot(epochs, val_acc, "b", label="Validation acc")
    plt.title(f"{model_name} - Training and Validation Accuracy")
    plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
    plt.legend()

    plt.tight_layout()
    plt.show()

# Plot results for each model
for model_name, results in results_history.items():
    plot_model_results(model_name, results['history'])
```





In conclusion, the model I would pick overall would be the model with MSE instead of binary\_crossentropy as the loss function. This model improved accuracy over all other models, and had a much lower loss function in both the training and validation sets.

## ✓ Outputting a table of all results together

```
import pandas as pd

# Convert the dictionary to a DataFrame
results_df = pd.DataFrame.from_dict(results_history, orient="index")

# Print the DataFrame as a table
print(results_df)

from tabulate import tabulate

# Print the results using tabulate
print(tabulate(results_df, headers='keys', tablefmt='pretty'))

# Extract accuracy and loss for each model from the results_history
accuracy = []
loss = []
models = list(results_history.keys())

for model_name, results in results_history.items():
    history_dict = results['history']
    accuracy.append(history_dict['accuracy'][-1]) # Use the last accuracy value
    loss.append(history_dict['loss'][-1]) # Use the last loss value
```



	loss	accuracy
base_model	0.294816	0.87980
model_with_added_layer	0.297208	0.88068
model_with_changed_hidden_units	0.298604	0.88480
model_with_mse	0.090495	0.88300
model_with_tanh	0.326985	0.86728
model_with_dropout	0.287226	0.88420

	history
base_model	{'accuracy': [0.7720666527748108, 0.8892666697...
model_with_added_layer	{'accuracy': [0.76746666431427, 0.891866683959...
model_with_changed_hidden_units	{'accuracy': [0.7624666690826416, 0.8868666887...
model_with_mse	{'accuracy': [0.7630666494369507, 0.8801333308...
model_with_tanh	{'accuracy': [0.787933349609375, 0.90113335847...
model_with_dropout	{'accuracy': [0.72079998254776, 0.850933313369...

	loss	accuracy
base_model	0.294816	0.87980
model_with_added_layer	0.297208	0.88068
model_with_changed_hidden_units	0.298604	0.88480
model_with_mse	0.090495	0.88300
model_with_tanh	0.326985	0.86728
model_with_dropout	0.287226	0.88420



	base_model		0.29481571912765503		0.879800021648407		{'accura
	model_with_added_layer		0.2972082197666168		0.8806800246238708		{'acc
	model_with_changed_hidden_units		0.29860374331474304		0.8848000168800354		{'accura
	model_with_mse		0.09049452841281891		0.8830000162124634		{'accura
	model_with_tanh		0.32698461413383484		0.8672800064086914		{'accu
	model_with_dropout		0.28722602128982544		0.8841999769210815		{'acc
+-----+-----+-----+-----+							

In conclusion, the model I would pick overall would be the model with MSE instead of binary\_crossentropy as the loss function. This model improved accuracy over all other models, and had a much lower loss function in both the training and validation sets.

## ✓ Plotting all results in charts to visually compare

```
# Plot accuracy comparison
plt.figure(figsize=(10, 5))
plt.bar(models, accuracy, color="green", label="Accuracy")
plt.ylabel("Accuracy")
plt.title("Model Accuracy Comparison")
plt.legend()
plt.xticks(rotation=45, ha="right") # Rotate model names for better readability
plt.show()

# Plot loss comparison
plt.figure(figsize=(10, 5))
plt.bar(models, loss, color="red", label="Loss")
plt.ylabel("Loss")
plt.title("Model Loss Comparison")
plt.legend()
plt.xticks(rotation=45, ha="right") # Rotate model names for better readability
plt.show()
```

